

Machine Learning vs. Deep Learning

Quan Khuc

July 11, 2026

Course: Machine Learning Fundamentals (AIML-500-01N1C) · **Assignment:** 2.3 · **Author:** Quan Khuc · **Date:** July 11, 2026

Introduction

Machine learning (ML) and deep learning (DL) are both subfields of artificial intelligence, but they solve problems in different ways. In traditional machine learning, a person decides in advance which characteristics of the data matter, extracts those features by hand, and feeds them to an algorithm such as logistic regression, a decision tree, or a support vector machine. In deep learning, a layered neural network learns the useful features on its own, directly from raw data, without a human naming them first.

Because of this difference, neither approach is universally better. The right choice depends on the structure of the problem: whether the signal can be captured by a short list of features a person can name, or whether it is buried in raw, high-dimensional data that no one can reduce to hand-written rules. The two examples below, both drawn from the lesson, show each case. Email spam filtering is well suited to traditional machine learning, and medical image diagnosis is well suited to deep learning. For each, I explain why the chosen approach fits and why the other approach does not.

Example 1 — Simple Machine Learning: Email Spam Filtering

Real-world application. Gmail uses machine learning to decide whether an incoming message is spam. Google reports that its filters block more than 99.9 percent of spam, phishing, and malware, keeping billions of unwanted messages out of inboxes every day. The same approach runs inside most corporate email gateways.

Why machine learning is the right fit. The signal that separates spam from legitimate mail lives in features a person can name and measure: the reputation of the sender's domain, whether the message fails authentication checks (SPF or DKIM), the presence of high-risk words such as "wire transfer" or "free prize," the number of links, the ratio of images to text, and the message length. Once these features are extracted, each email becomes a short, structured row of numbers. Algorithms such as Naive Bayes and logistic regression model the relationship between those features and the spam label accurately, and they do so cheaply. Training needs only modest hardware and a moderate amount of labeled mail, and the model can be retrained in minutes. That speed matters because spammers change their tactics constantly, and the filter has to keep up. The approach is also transparent: the model can point to the specific features that caused it to flag a message, which is important when a legitimate business email is wrongly quarantined and a person has to review why.

Why deep learning is not suitable. Deep learning earns its advantage when features must be discovered automatically from raw, unstructured data. Spam filtering does not have that problem, because the informative features are already known and easy to compute. A deep neural network would add real costs without a matching benefit. It would require a much larger labeled dataset and specialized hardware such as GPUs to train, yet it would not classify email meaningfully better than a well-tuned linear model. It would also be slower to run at the scale of billions of messages a day, and harder to retrain quickly as spam evolves. Finally, a deep network is difficult to explain, which is a liability when a user disputes a blocked message or an administrator needs to justify the decision. For this problem, deep learning is more machinery than the task requires.

Example 2 — Deep Learning: Medical Image Diagnosis

Real-world application. Deep learning is used to read medical images and flag disease. Google Health built a system that detects diabetic retinopathy, a leading cause of blindness, from photographs of the retina. Trained on a large set of expert-graded images, the model matched the accuracy of board-certified ophthalmologists, and it has been deployed to screen patients in clinics in India and Thailand where specialists are scarce.

Why deep learning is the right fit. A retinal photograph is made of millions of pixels, and the evidence of disease, tiny hemorrhages and microaneurysms, is a subtle spatial pattern that no one can reduce to a short list of hand-written rules. This is exactly the situation deep learning was built for. A convolutional neural network learns features in layers: early layers detect simple edges and textures, and deeper layers combine them into the lesions and structures that indicate disease. The network discovers these features on its own, directly from the images, so no human has to specify them in advance. Convolutional networks also share weights across the image, so a lesion is recognized no matter where it appears, and they handle the natural variation in lighting, angle, and anatomy across real patients. Given enough labeled images, this approach reaches specialist-level accuracy, and its performance improves as more data becomes available.

Why traditional machine learning is not suitable. Traditional machine learning depends on a person choosing the features first, and for medical images no one can. To use a support vector machine or a decision tree, an engineer would have to hand-craft visual descriptors such as edge detectors or color histograms, but these generic features are too crude to capture the specific patterns that distinguish grades of disease. Feeding the raw pixels straight into a traditional algorithm fails just as badly: the input is far too high-dimensional, and the relationship between individual pixels and a diagnosis is far too complex and non-linear for these models to learn.

The core obstacle is that the diagnostic pattern is too intricate for any human to specify by hand, and specifying features by hand is precisely what traditional machine learning requires. That gap is what deep learning closes.

Conclusion

These two examples mirror each other and make the same point from opposite directions. Spam filtering succeeds with traditional machine learning because its signal is captured well by a handful of features a person can name, which makes deep learning unnecessary overhead. Medical image diagnosis requires deep learning because its signal is a complex spatial pattern hidden in raw pixels that no person can reduce to hand-written features, which is exactly where traditional machine learning breaks down. The lesson is not that one method is superior, but that the method should match the problem: use traditional machine learning when the important features are known and easy to extract, and use deep learning when the features must be learned from raw, high-dimensional data.

References

- Machine Learning Fundamentals (AIML-500), Module 2 lesson: “Machine Learning vs. Deep Learning.”
- Gulshan, V., et al. (2016). “Development and Validation of a Deep Learning Algorithm for Detection of Diabetic Retinopathy in Retinal Fundus Photographs.” *JAMA*, 316(22), 2402–2410.
- Google Safety Center and Google Workspace, published statistics on Gmail spam and phishing filtering.